Optimisation of d2q9-bgk Lattice Boltzmann Scheme

James Elgar, za18968

March 5, 2021

1 Introduction

This report will explore the optimizations applied to the implementation of a given algorithm to solve a d2q9-bgk Lattice Boltzmann scheme. Starting with an unoptimized serial implementation, it will apply three stages of optimisation and compare the results. The first stage will be serial optimisations which make the general alogirthm more efficient mainly by reducing the number of parses over the cells array to one, reducing the memory bandwidth required. The second stage will make use of SIMD, single instruction multiple data, where the compiler can vectorize certain sections of the algorithm allowing it to compute multiple data values with single instructions by making use of vector registers and instructions. The final stage will explore using OpenMP to run section of the algorithm in parallel and convey the large benifits of being able to spread out the required computation over multiple cores. However it will also look at the limitation of this approach as memory bandwidth and sharing can continue to limit performance. Throughout the stages few changes will be made to the underlying implementation of the algorithm and the majority of the optimisation will come from assisting the compiler and the OpenMP library with the optimisations they are able to achieve.

2 Serial Optimisation

2.1 Compiler options

Many modern compilers provide various options to optimize the given code, simply by adding compiler flag options. The "O" flags offer various levels of optimisation with `-O3` being the highest.

Table 1 shows the run time of the serial optimized code (without vectorization) with different compiler flags. This shows that by simply using the `-Ofast` flag, when compared to no optimisations (`-O0`), there can be a 4x decrease in run time. In later sections additional compiler flags were added to optimize the other stages (such as vectorization). The compiler flags used in the final implementation were

- `-fast` - icc specific flag which includes -Ofast and a few other optimisations

- `-mtune=native` and `-xAVX` allow for the compiled binary to be specific to the current environment. `xAVX` explicitly specific the type of SIMD instructions to target (based on the architecture).

- `-no-prec-sqrt` reduces the required precision of the square root function.

Results	
Compiler flag	Run time
-O0	111.7
-O1	34.5
-O2	34.2
-O3	31.9
-Ofast	26.8
-Ofast -mtune=native	26.5

Table 1: Serial run time with different compiler flags

2.2 Reduce Memory Accesses

In the original implementation of the algorithm, for each timestep the `cells` array was looped over 4 times, in `propagate`, `rebound`, `collision` and `av_velocity`. This resulted in repeated stores and loads of the same sections of memory. To prevent this the first 3 separate functions (those in the `timestep` function) were fuzed into a single loop. This meant the same sections of memory were used closer together making it is more likely for them to still be in cache. This therefore reduces the memory bandwidth used and since this implementation is memory bound results in a increase in performance. The function `av_velocity` was also repeating calculations that already took place in collision and requiring an additional loop over cells. The result was therefore calculated in the single parse over the cells, reducing the need for this additional loop.

Another issue with the original implementation was in propagate and collisions where the values were switched between `tmp_cells` and `cells` multiple times which results in a waste of memory bandwidth as unnecessary loads and stores take place. In the new implementation the `tmp_cells` array was used as the "answer" space and stores only the next timestep's cell values. This then only required a single write to `tmp_cells` each timestep. At the end of the timestep the `tmp_cells` and `cells` array's pointers were then swapped which set the `cells` array to the correct value without having to write directly to the array. Overall these optimisations reduced the runtime of the serial implementation from 38.5 seconds to 26.4.

2.3 Vecotrisation

Having improved the implementation of the serial code, there is now a clear critical section (the section which takes the most time), inside the single pass of the cell's grid. This is where the majority of the computation takes places and thus is where most of the time of the program is used. Since this section preform repeat operations on the input array, there is an additional serial approach to optimization.

Modern CPU architectures include vector registers and instructions. This allows for computation to take place on a vector of values whilst only using a single instruction. This approach is known as SIMD, (single-instruction multiple data). When enabled, compilers are able to automatically make use of these vector instructions. For this algorithm, this has the

potential for large performance gains as fewer instructions are required for the cells array in the critical section.

```
1   // Allocate aligned data
2   cells_ptr->speed0 = _mm_malloc(params->nx * params
      ->ny * sizeof(float), 64);
3   // When using the array allow the compiler to
      assume alignment
4   __assume_aligned(cells->speed0, 64);
```

Listing 1: Example of memory allignment for a cells array.

In many cases the compiler is able to automatically vectorize code blocks. However if the data is not aligned then the compiler will have to complete a prior step, called the "prologue"[1] step, to process this unaligned data, or alternatively use unaligned instructions (which are less efficient). The compiler is also not able to assume that an array of floats, such as the arrays of speeds in the new `t_speed` struct are aligned. Therefore it will have to complete this prior step regardless. To prevent this the arrays can be aligned when they are created using the `__mm_malloc` function, as shown in Listing 1. This ensure the created array is aligned on the request boundry, in this case 64 bytes was used to match the size of the cache line in an *intel xeon e5-2680*. Compiler directives can then be use to tell the compiler that the arrays are aligned. This allows the compiler to skip the "proluge" step and use the aligned instructions. This significantly reduces the overhead of vectorizing as well as making it more likely for the compiler to automatically vectorize the given section.

Another useful hint to assist the compiler in its optimizations is to use the `restrict` keyword. This tells the compiler to prevent pointer aliasing, which ensures the target value will only be accessed through the restricted pointer (not another aliased pointer). In practical terms this reduces the number of times the pointer value has to be fetched from memory as the cached value can be used repeatedly without the risk of it having changed as a result of an overlap between pointer (in this case there would be a risk of overlap with `cells` and `tmp_cells`). In a memory bound problem this can therefore be an important step in reducing the memory bandwidth used.

```
1   __assume(params.nx%128==0);
2   __assume(params.ny%128==0);
```

Listing 2: Additional compiler hints

The intel compiler also allows for other hints to be provided. In this implementation hints were added to inform the compiler about the size of the cells array (and therefore the size of the loops). As show in the example in Listing 2, the `__assume` directive was used, to tell compiler that the size of the cells array was divisible by various given powers of 2. In the implementation all powers of 2, up to a maximum of 128 were provided (as this was the maximum input size). This should further assist the compiler in its optimizations, especially when vectorizing the inner loop.

The final step to ensuring the code was vectorized was to add the compiler directive `#pragma omp simd` to the inner for loop. This explicitly tells the compiler that this section of the code should be vectorized.

This alone, although successfully vectorizing the loop, did not result in a significant performance gain. Observing the roofline graph showed that this version of the vectorized code was very heavily memory bandwidth bound. This was as a result of the data layout of the cells array.

In the original implementation the speeds for each cell were stored in a structure (`t_speed`) and the whole grid of cells was an array of these structures. This implementation does not lend it self to vetorization as it can result in unnecessary fetches from memory. This is because when a single speed value is fetch an entire cache line (64 bits) will be used to fetch the structure. This results in a waste of memory bandwidth as the rest of the cache line is not used and only the single float value for that speed is actually required. Therefore switching the implementation to use a structure of arrays (where each speed is a separate array) allowed the code to be vectorized more efficiently as each fetch for a speed only fetched a single float and therefore reduced the overall memory bandwidth required.

When comparing the roofline results in Figure 1 we can see a large improvement in the performance, from 2.49 GLOPs to 13.2 GLOPs, once the vectorization is introduced. Similarly we also see a reduction in the run times as shown in Table 3. These improvements are as a direct result of vectorization and the update to the data layout which reduces the memory bandwidth requirements. The intel compiler also provides a vectorized implementation for the `sqrt` function which allows for the majority of the computation in the critical section to be vectorized.

Despite the large improvements, it may have been expected to see a larger improvement when the majority of instructions are now vectorized, which in the critical section could reduce the number of required instruction by a significant factor since the vector length is 32 bits. However such a large improvement is not observed as the problem remains memory bandwidth bound meaning as the number of instruction reduces and the size of the vectors increase, the memory bandwidth required increase as more data is required at the same time. Moreover there is some additional overhead to vectorizing the code and thus this will impact the overall run time.
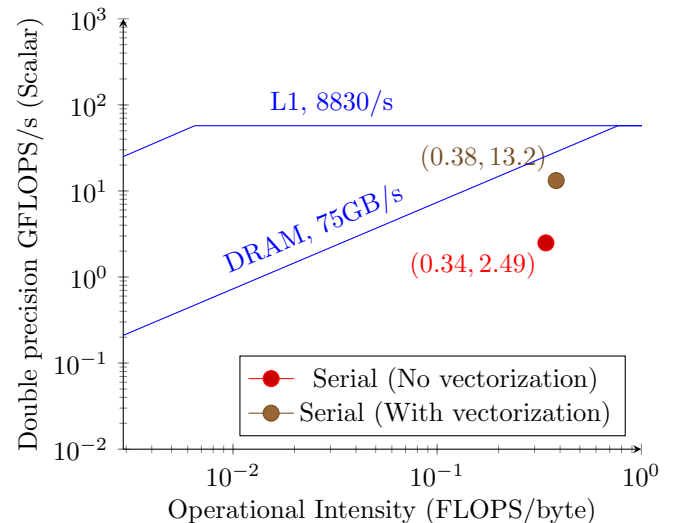

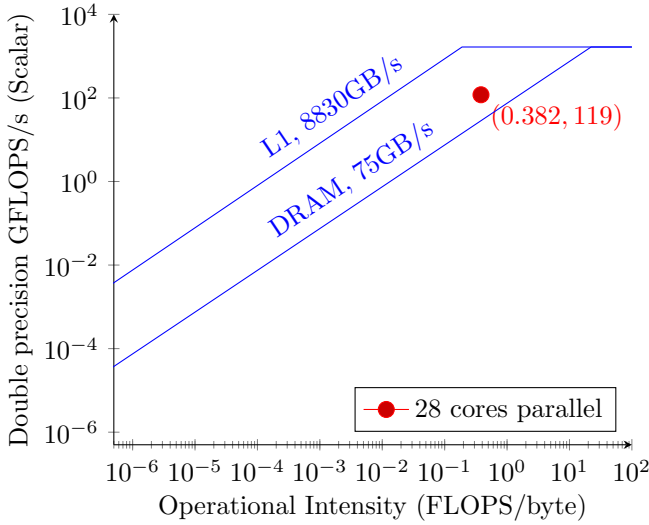
Figure 1: Roofline analysis for serial implementations

Figure 2: Roofline analysis for parallel implementations

3 Parallel (OpenMP)

Having optimized the inner loop in the above section the next optimization that can be used is to run the loops in parallel. In this implementation the outer loop was run in parallel while the inner loop maintained the other optimizations. This parallelism was achieved by using the OpenMP library. This library provides a collection of compiler directives which can be used to write parallel shared memory programs.

The first step in order to run the loop in parallel is to turn on OpenMP for the compiler with the `-qopenmp` flag. With this the OpenMP library is available to use. At a basic level adding the `parallel` and `for` constructs should allow the loop to run in parallel. The `parallel` directive tells the compiler to run the preceding block in parallel and `for` tells it to distribute the work of the following for loop evenly across the threads.

It is then important to handle the memory sharing that takes place in the loop. Since the `av_velocity` is also calculated in this loop each loop is required to calculate the total number of cells without obstacles and the sum of the magnitude of the velocity for these cells. When running the code with just the `#pragma omp parallel for` directive, the wrong answer is obtained. This is because the values required to calculate `av_velocity` are in shared memory and thus every thread is competing to set the shared values rather than accumulating the values across all the threads.

In order to efficiently calculate these shared sums a `reduction` is required. The `reduction` directive tells the compiler that the given variable is going to be accumulated over the duration of the parallel region. This allows each thread to calculate its own sum without having to syncronise the global shared value on every write. Once the forked section is complete and the threads syncronise then all the sums can be combined to determine the final value of the accumulated variables. This is signifantly faster than requiring a syncronised section during the fork as the threads can run independently and not have to wait for each thread to syncronise the sum or lock and unlock (in the case of an atomic variable) the values.

Another consideration when writing a parallel program is the number of cores to use. In a bluecrystal node there are 2 sock-

ets, which house 2 *intel xeon e5-2680* CPUs. Each of these CPUs have 14 cores, giving a total of 28 cores per node. Table 2 and Figure 3 show the run times of the parallel code on different numbers of cores. In the larger problems (1024x1024) it is clear that as the number of cores increases the run time decreases. Initially, between 1 and 4 cores, the rate of decrease is linear, meaning as the number of cores increases the run time is halved. This is as a result of the work in the critical section being shared over a larger pool of cores and thus the total time of the forked region is decreased.

However this linear relationship becomes a sublinear plateaus as the number of cores increases. The reason for this can be observed in Figure 2 which shows the roofline analysis for the parallel algorithm running on 28 cores. The results lie in the memory bound region, this is as a result of the memory bandwidth becoming saturated as the number of cores are increased and thus the computation rate increases at a faster rate than the memory bandwidth available.

However in the smaller problems there is a reduction in run time after 14 cores. This is observed as once the core count is above 14, two sockets are in use (both CPUs) and therefore memory sharing between the sockets is required. This is where a shared section of memory is in the cache of the other socket so in order to access that section of memory a request has to be routed through the other socket. This takes approximately two times as long as fetching from the socket's local cache. For this reason when running the algorithm on smaller grid sizes (such as 128x128) only 14 cores were used. This removes this problem completely as only a single socket is used so all cache accesses are for the local CPU's cache.

In general this non-uniform memory access, NUMA, can be mitigated by pinning the threads to a specific core. Setting `OMP_PROC_BIND` to `close` tells OpenMP to bind the threads to specific cores and also as the threads are assigned, ensure they are assigned to the threads closest to the master (main) thread. Assigning `OMP_PLACES` to `cores` ensures that each core is used as a single place (and no core sharing occurs such as hyper-threading). Although the threads are now pinned to individual cores this does not mean that the NUMA problem is reduced. This is because during the memory assignment, in the `initialise` function, the initial values for the cells array are assigned serially (on a single thread). This means all the values will in exist the same cache, linked to this single NUMA region. This can be prevented by ensuring the initialisation of the values is parallelised in the same way as the it is accessed in the critical region. In this case this means also using the `#pragma omp parallel for` directive when assigning the initial values to the cells array. This will ensure the data is allocated in the same NUMA region to which it is going to be required in the critical loop. The combination of these factors should therefore reduce the number of socket-to-socket memory requests allowing for the average time for each memory access to be reduced.

When comparing the run times of the parallel code with the vectorized implementation in Table 3, for larger problem sizes we see an almost 14x reduction. Similarly when comparing the roofline analysis in Figure 1 and Figure 2 an over 9 times improvement in performance is observed. However similarly to the vectorized code we do not observe a 28 times improvement as may be expected when running over 28 cores. This

is as a result of the memory saturation and the overhead of forking and then syncronising the code. In smaller problems this overhead and the need for memory sharing can even result in reduced performance as the core count increases after a certain point.
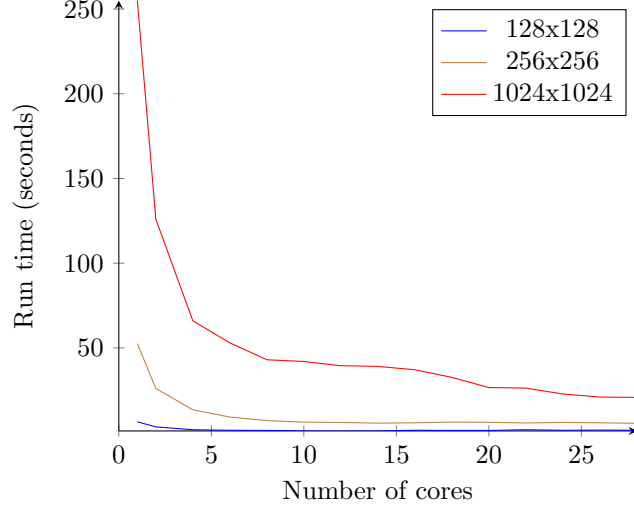

Figure 3: Parallel scaling results for 1-28 cores on 3 grid sizes

Results			
Number of cores	128x128	256x256	1024x1024
1	6.4	52.5	255.4
2	3.4	26.1	126.2
4	1.7	13.5	66.1
6	1.4	9.2	53.3
8	1.3	7.1	43.4
10	1.0	6.2	42.4
12	0.9	6.0	39.7
14	1.0	5.6	39.1
16	1.4	5.8	37.1
18	1.3	6.2	32.6
20	1.3	6.1	26.6
22	1.7	5.7	26.3
24	1.4	6.0	22.8
26	1.5	5.8	21.0
28	1.4	5.5	20.8

Table 2: Parallel scaling results for 1-28 cores

Results			
Grid size	Serial	Vectorised	Parallel(cores)
128x128	26.4	7.4	1.15(14)
128x256	53.3	16.1	1.85(14)
256x256	213.6	59.6	5.68(28)
1024x1024	869.1	318.9	20.8(28)

Table 3: Run times for optimisation stages

4 Conclusion

Modern compilers offer many powerful options to assist in code optimisation. Often the most efficient approach to optimising a code can be to assist the compiler in its own optimisations. Simple changes such as using compiler optimisation flags can have a profound impact on the runtime of the algorithm. Similarly with the correct hints and memory aligned compilers are able to automatically vertorize a given code. This can dramatically reduce the run time by allowing a vector instructions to compute multiple values and thus reduce the overall number of instructions required. In this implementation, the vectorized code ran over 2.5 times faster than the optimised serial version with an approximately 5 times increase in performance. Once a code is well optimised in serial further improvements can be gained by making sections of the program run in parallel. This can dramatically increase the performance of the code as the work is shared across multiple cores. With modern processors often having high core counts, this can also result in large performance gains. Similarly to vectorization, using OpenMP directives and environment variables, parallelized code can be achieved with little changed to the implementation. In this running the critical sections in parallel resulted in up to a 13 times reduction in overall run time on the largest problem size.

References

[1] Kevin O'Brien Alexandre E. Eichenberger, Peng Wu. Vectorization for simd architectures with alignment constraints.